

NAYANA DAVIS

SQL JOINS

WHAT WE'VE LEARNED TO DO WITH DATABASES SO FAR

- ▶ how to connect to a local or remote db
- ▶ how to add, remove, edit data
- ▶ how to perform simple queries
- ▶ how to aggregate, group and sort data

SQL COMMANDS

- ▶ SELECT
- ▶ CREATE
- ▶ INSERT
- ▶ DELETE
- ▶ UPDATE
- ▶ ORDER BY
- ▶ HAVING
- ▶ LIKE
- ▶ DISTINCT
- ▶ LIMIT

NORMALIZED VS DENORMALIZED DATA

- ▶ Normalized structures have a single table per entity, and use many foreign keys or link tables to connect the entities.
- ▶ Denormalized tables have fewer tables and may (for example) place all of the tweets and the information on users in one table.

NORMALIZED VS DENORMALIZED DATA

| Normalized | | | | | |
|--|-------------------|---------|--------|---------------|---------|
| Normalized – Data is broken into multiple tables | | | | | |
| Product | | Color | | Product-Color | |
| ProductID | Desc | ColorID | Desc | ProductID | ColorID |
| 1 | Mtn Bike #778 | 1 | Red | 1 | 1 |
| 2 | Road Bike #123 | 2 | Black | 1 | 2 |
| 3 | Touring Bike #222 | 3 | Silver | 2 | 1 |
| | | 4 | Mauve | 2 | 2 |
| | | | | 2 | 3 |
| | | | | 3 | 1 |
| | | | | 3 | 3 |
| | | | | 3 | 4 |


| Denormalized | | | | |
|------------------------------|----------------|---------|-------------------|--------|
| Denormalized – Data combined | | | | |
| Product | (denormalized) | | | |
| ProductSK | ProductID | ColorID | Desc | Color |
| 1 | 1 | 1 | Mtn Bike #778 | Red |
| 2 | 1 | 2 | Mtn Bike #778 | Black |
| 3 | 2 | 1 | Road Bike #123 | Red |
| 4 | 2 | 2 | Road Bike #123 | Black |
| 5 | 2 | 3 | Road Bike #123 | Silver |
| 6 | 3 | 1 | Touring Bike #222 | Red |
| 7 | 3 | 3 | Touring Bike #222 | Silver |
| 8 | 3 | 4 | Touring Bike #222 | Mauve |

NORMALIZED VS DENORMALIZED DATA

Normalized tables save the storage space by separating the information. However, if we ever need to access those two pieces of information, we would need to join the two tables, which can be a fairly slow operation.

JOINS

The natural way to merge data in SQL



JOINS

- ▶ SQL joins are used when data is spread in different tables. A join operation allows to combine rows from two or more tables in a single new table. In order for this to be possible, a common field between the tables need to exist.
- ▶ Join operations can be thought of as operations between two sets, where records with the same key are combined and records missing in one set are either discarded or included as NULL values.


FOREIGN KEY VS PRIMARY KEY

| Primary Key | Foreign Key |
|---|--|
| Primary key uniquely identify a record in the table. | Foreign key is a field in the table that is primary key in another table. |
| Primary Key can't accept null values. | Foreign key can accept multiple null value. |
| By default, Primary key is clustered index and data in the database table is physically organized in the sequence of clustered index. | Foreign key do not automatically create an index, clustered or non-clustered. You can manually create an index on foreign key. |
| We can have only one Primary key in a table. | We can have more than one foreign key in a table. |

Note: clustered index means it impacts the way records are stored in a table


INNER JOIN

This is the simplest, most understood Join and is the most common. This query will return all of the records in the left table (table A) that have a matching record in the right table (table B).




LEFT JOIN

This query will return all of the records in the left table (table A) regardless if any of those records have a match in the right table (table B). It will also return any matching records from the right table.




RIGHT JOIN

This query will return all of the records in the right table (table B) regardless if any of those records have a match in the left table (table A). It will also return any matching records from the left table.



FULL JOIN/OUTER JOIN/FULL OUTER JOIN

This Join can also be referred to as a FULL OUTER JOIN or a FULL JOIN. This query will return all of the records from both tables, joining records from the left table (table A) that match records from the right table (table B)



DEMO: INNER JOIN

Let's connect to the Northwind database

```
psql -h dsi.c20gkj5cvu3l.us-east-1.rds.amazonaws.com -p  
5432 -U dsi_student northwind
```

password: gastudents

DEMO: INNER JOIN

Let's look at all our tables in the database and then check out the orders and customers tables in full

\dt

TABLE customers;

TABLE orders;

DEMO: INNER JOIN

Let's consider a few columns of the orders table:

| OrderID | CustomerID | OrderDate |
|---------|------------|------------|
| 10308 | 2 | 1996-09-18 |
| 10309 | 37 | 1996-09-19 |
| 10310 | 77 | 1996-09-20 |

And a few in customers:

| CustomerID | CompanyName | ContactName | Country |
|------------|------------------------------------|----------------|---------|
| 1 | Alfreds Futterkiste | Maria Anders | Germany |
| 2 | Ana Trujillo Emparedados y helados | Ana Trujillo | Mexico |
| 3 | Antonio Moreno Taquería | Antonio Moreno | Mexico |

DEMO: INNER JOIN

Notice that the CustomerID column in the Orders table refers to the CustomerID in the Customers table. The relationship between the two tables from before is the CustomerID column. We can thus JOIN the two tables in order to obtain a table like the following, where the information contained in the two tables is joined in a single table, using the common key of CustomerID.

| OrderID | CompanyName | OrderDate |
|---------|------------------------------------|------------|
| 10308 | Ana Trujillo Emparedados y helados | 9/18/1996 |
| 10365 | Antonio Moreno Taquería | 11/27/1996 |
| 10383 | Around the Horn | 12/16/1996 |
| 10355 | Around the Horn | 11/15/1996 |
| 10278 | Berglunds snabbköp | 8/12/1996 |

DEMO: INNER JOIN

```
SELECT <select_list>  
FROM Table_A A  
INNER JOIN Table_B B  
ON A.Key = B.Key
```

```
SELECT orders."OrderID", customers."CompanyName", orders."OrderDate"  
FROM orders  
INNER JOIN customers  
ON orders."CustomerID"=customers."CustomerID";
```

YOU DO: LEFT JOIN (3 MINUTES)

Syntax:

```
SELECT column_name(s)
FROM table1
LEFT JOIN table2
ON table1.column_name=table2.column_name;
```

Perform a left join with orders as table 1, customers as table 2

Slack query in students channel

YOU DO: RIGHT JOIN (3 MINUTES)

Syntax:

```
SELECT column_name(s)
FROM table1
RIGHT JOIN table2
ON table1.column_name=table2.column_name;
```

Perform a right join with orders as table 1, customers as table 2

Slack query in students channel

YOU DO: FULL JOIN (3 MINUTES)

Syntax:

```
SELECT column_name(s)
FROM table1
FULL OUTER JOIN table2
ON table1.column_name=table2.column_name;
```

Perform a full join with orders as table 1, customers as table 2

Slack query in students channel

YOU DO: WRITE QUERIES TO ANSWER THESE QUESTIONS(15 MINUTES)

Look over what you previously learned about SQL:

<https://github.com/ga-students/DSI-DC-1/tree/master/week-05>

QUESTIONS

Answer these questions:

How many products per category does the catalog contain? Print the answer with the `CategoryName`, and `Count`.

What 5 customers are generating the highest revenue? Print a table with `CustomerID` and `Total Revenue`. You will need to use data from 3 tables.

In which country are the top 5 suppliers by number of units supplied? Print a table with the supplier's `CompanyName`, `Country` and total number of units supplied.

Slack your code in the students channel

SUBQUERIES

- ▶ A Subquery or Inner query or Nested query is a query within another SQL query. It is used to further restrict the data to be retrieved by returning data that will be used in the main query as a condition.
- ▶ Subqueries can be used with the SELECT, INSERT, UPDATE, and DELETE statements along with the operators like =, <, >, >=, <=, IN, BETWEEN etc.

YOU DO: SUBQUERIES (5 MINUTES)

Syntax:

```
SELECT column_name1
  FROM table_name1
 WHERE column_name2 [Comparison Operator]
      (SELECT column_name3
        FROM table_name2
        WHERE condition);
```

Perform a query that uses a subquery that extract all the orders from customers based in France.

Do this again, but this time use a JOIN operation

INDEPENDENT PRACTICE (10 MINUTES)

- ▶ Work in pairs: go to <http://www.w3schools.com/sql> and choose a command you have not hear of. Read about it for 5 minutes, then explain it to your pair (take 2.5 minutes turns).
- ▶ Use the last 5 minutes to share some interesting finds with the rest of the class.